

# T4\_MNIST\_Nihal\_JG

March 31, 2020

This aim of this Problem Statement is to introduce Deep Learning, which is one of the most used techniques for Computer Vision Application. Before beginning go thorough the following articles: 1. <https://www.nature.com/articles/nature14539> : It is fairly easy to follow and provides an excellent overview of the field. (You will need to login with gmail to download the pdf.)

2. <http://deeplearning.stanford.edu/tutorial/supervised/MultiLayerNeuralNetworks/> : A short introduction to Multi-layer perceptrons.
3. <https://cs231n.github.io/convolutional-networks/> : Introduction to Convolutional Neural Networks. CNNs are generally used for computer vision problems.

This introduction should be sufficient to get you started with this problem statement. If after going through the PS, you guys are interested in further exploring the field, I would suggest the following resources: 1. <http://cs231n.stanford.edu/> : Online course by Stanford. 2. <http://introtodeeplearning.com/>: Video Lectures from MIT 3. <https://www.deeplearningbook.org/>: Most popular book on Deep Learning

The aim of this get you all familiarized with Deep Learning in PyTorch, a very popular Deep Learning Library (or in general GPU computation library).

Some of the preprocessing work has been done, and you guys are expected to fill-in code where you are asked to.

In this notebook we will be doing the following tasks: 1. Define a small convolutional neural network and train in on MNIST digit dataset (as a classification task). 2. (Bonus) Take the trained network, freeze all of its weights and learn the image which gives the output of a particular digit. (This will be defined better, when we reach that task). 3. (Bonus) Try the same on a different dataset (like MNIST fashion dataset). 4. (Bonus) Try to formulate the MNIST classification problem as a regression problem instead. (i.e. Have 1 output unit which outputs a floating point value and you round it off to obtain the digit.)

Before diving into the code ensure that you copy the notebook to your drive (See the option in File Tab) and that the Runtime Type is set to GPU (Runtime tab -> Change runtime type). To see the importance of GPU in deep learning see [this](#) short article.

The following cell imports all the necessary packages. We will be using : 1. Several Modules of PyTorch. This will be used to do all the processing and importing the dataset etc. Read [this](#) and [this](#) article to get started. Here are the [official docs](#).

2. NumPy is likely the most popular Matrix manipulation and linear algebra library available for Python. Unfortunately it does not support processing on GPU and therefore we will not be using it much here (Only as a support library for matplotlib). Here is a quick [tutorial](#) if you are interested.

3. Matplotlib is the most popular library for generating plots in Python. We will be using it for printing images.

```
[1]: import torch
import torchvision
import torch.nn as nn
import torch.nn.functional as F
import torch.utils.data
from torch.autograd import Variable

import numpy as np

from matplotlib import pyplot as plt

#This to ensure that the default device for processing PyTorch tensors is GPU
→instead of a CPU.
device = 'cuda' #Change this to 'cpu' to run on cpu.
```

Here we will download the MNIST digit dataset and convert it into PyTorch dataset object.

torchvision module of PyTorch provides several Computer Vision specific functionalities one of which is easy importing of several [popular datasets](#).

```
[2]: #We will load the MNIST dataset here.

#This creates a transform object which will be passed to the later functions.
→It converts the training image from PIL Image objects to PyTorch image
→objects.
transform = torchvision.transforms.Compose([torchvision.transforms.ToTensor()])

train_data = torchvision.datasets.MNIST(root = './data', train = True, download_
→= True, transform = transform )

test_data = torchvision.datasets.MNIST(root = './data',train = False ,download_
→= True, transform = transform)

print(train_data)
print(test_data)

plt.imshow(np.asarray(train_data[1][0].reshape(28,28)))
```

Dataset MNIST


Number of datapoints: 60000  
Root location: ./data  
Split: Train

```

StandardTransform
Transform: Compose(
  ToTensor()
)
Dataset MNIST
  Number of datapoints: 10000
  Root location: ./data
  Split: Test
  StandardTransform
Transform: Compose(
  ToTensor()
)

```

[2]: <matplotlib.image.AxesImage at 0x7f6f0d3d4d30>



The following cell converts the DataSet objects to DataLoader objects. DataLoader objects in PyTorch provide several useful functionalities, such as automatic batching, parallel data processing on the CPU, etc. Therefore it is generally advisable to use dataloader instead of manually handling data.

Read more about PyTorch DataSet and DataLoader [here](#)

```

[3]: #BATCHs are only used for training, while testing we conventionally use a batch
    ↪size of 1.
    BATCH_SIZE = 32
    train_loader = torch.utils.data.DataLoader (train_data, batch_size =
    ↪BATCH_SIZE, )

```

```
test_loader = torch.utils.data.DataLoader (test_data, batch_size = 1, )

print(train_loader)
```

<torch.utils.data.dataloader.DataLoader object at 0x7f6f0cce9f60>

## TASK 1

We will define our network in the following cell. There are several ways of doing this in PyTorch, but in this example we will do it by subclassing `nn.Module`. You should be able to do this after reading the tutorials given at the top of the page. [This](#) is extra reading to help you out. (Do remember that we are dealing with 1 channel B/W images not 3 channels colour images.)

[Here](#) are official docs of `nn` module to help you out.

Our network will have 4 convolutional layers followed by 2 fully-connected layers. (You free to change make slight changes to this but don't make the network too deep in the final submission).

Don't forget to have ReLU activations and maxpooling layers.

Experiment with different kernel sizes, width and sizes of hidden fully connected units.

```
[4]: ## Complete the functions __init__ and forward below

class MyConvNet(nn.Module):
    def __init__(self):
        """
            Initialize all the weights for each layer you require.
        """

        super(MyConvNet, self).__init__()

        # __begin

        self.conv1 = nn.Conv2d(1, 16, kernel_size=5)
        self.conv2 = nn.Conv2d(16, 32, kernel_size=5)
        self.conv3 = nn.Conv2d(32, 64, kernel_size=3)
        self.conv4 = nn.Conv2d(64, 64, kernel_size=3)

        self.fc1 = nn.Linear(8 * 8 * 4, 256)
        self.fc2 = nn.Linear(256, 10)
        # __end

    def forward (self, x):
        """
            Define a forward pass for the given network. x is a 4 dimensional input_
            ↪ tensor.
            Input: x.size() = [BATCH_SIZE, 1, 28, 28]
```

Return a 2-dimensional, 10-unit output tensor representing the  
 →probabiCalculated padded input size per channel: (4 x 4). Kernel size: (5 x  
 →5). Kernel size can't be greater than actual input sizelity of each class.  
 →Use a softmax output unit.

```

Output: x.size() = [BATCH_SIZE, 10]
"""

# Put your code here:
# __begin

#print(x.shape)
x = F.relu(self.conv1(x))

#x = F.dropout(x, p=0.5, training=self.training)
x = F.relu(F.max_pool2d(self.conv2(x), 2))
#print(x.shape)
#x = F.dropout(x, p=0.5, training=self.training)

x = F.relu(F.max_pool2d(self.conv3(x), 2))
#print(x.shape)
#x = F.dropout(x, p=0.5, training=self.training)

x = F.relu(F.max_pool2d(self.conv4(x), 1))
#print(x.shape)
#x = F.dropout(x, p=0.5, training=self.training)

x = x.view(-1, 8*8*4 )
#print(x.shape)
x = F.relu(self.fc1(x))
#print(x.shape)
#x = F.dropout(x, training=self.training)

x = self.fc2(x)
#print(x.shape)
# __end
#Return statement uses softmax, so after the final fully connected layer
→store the value in x (without using ReLU or any other activation).
return F.softmax(x, dim = 1)

```

The following cell will instantiate the the network we defined above. Check the size of the output tensor (it should be [BATCH\_SIZE, 10]) and see if the values in the output tensor printed below are numbers between 0 and 1. (Infact most values should be close to 0.1. Why?)

[5]: #Setting few things up.

```

net = MyConvNet()
net.cuda()
net.to(device)

```

```

print(net)
tld = iter(train_loader)
im = next(tld)[0].to(device)
print('Size of the output tensor:', net.forward(im).size())
print(net.forward(im))

```

```

MyConvNet(
  (conv1): Conv2d(1, 16, kernel_size=(5, 5), stride=(1, 1))
  (conv2): Conv2d(16, 32, kernel_size=(5, 5), stride=(1, 1))
  (conv3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
  (conv4): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1))
  (fc1): Linear(in_features=256, out_features=256, bias=True)
  (fc2): Linear(in_features=256, out_features=10, bias=True)
)
Size of the output tensor: torch.Size([32, 10])
tensor([[0.0964, 0.0995, 0.1016, 0.1015, 0.0954, 0.0982, 0.0980, 0.1004, 0.1060,
        0.1029],
        [0.0967, 0.0992, 0.1017, 0.1016, 0.0954, 0.0983, 0.0975, 0.1005, 0.1061,
        0.1030],
        [0.0968, 0.0994, 0.1016, 0.1014, 0.0952, 0.0980, 0.0979, 0.1005, 0.1064,
        0.1028],
        [0.0965, 0.0993, 0.1018, 0.1016, 0.0956, 0.0976, 0.0979, 0.1006, 0.1063,
        0.1028],
        [0.0966, 0.0995, 0.1019, 0.1012, 0.0954, 0.0981, 0.0980, 0.1002, 0.1064,
        0.1027],
        [0.0967, 0.0991, 0.1016, 0.1012, 0.0957, 0.0977, 0.0979, 0.1007, 0.1065,
        0.1029],
        [0.0965, 0.0999, 0.1017, 0.1011, 0.0956, 0.0983, 0.0978, 0.1004, 0.1057,
        0.1031],
        [0.0966, 0.0994, 0.1014, 0.1018, 0.0954, 0.0975, 0.0978, 0.1007, 0.1064,
        0.1029],
        [0.0967, 0.1000, 0.1020, 0.1010, 0.0953, 0.0984, 0.0981, 0.1003, 0.1055,
        0.1029],
        [0.0968, 0.0990, 0.1016, 0.1016, 0.0956, 0.0976, 0.0980, 0.1003, 0.1064,
        0.1031],
        [0.0964, 0.0994, 0.1017, 0.1013, 0.0955, 0.0978, 0.0978, 0.1010, 0.1063,
        0.1027],
        [0.0966, 0.0995, 0.1015, 0.1015, 0.0957, 0.0979, 0.0979, 0.1005, 0.1063,
        0.1026],
        [0.0964, 0.0994, 0.1013, 0.1016, 0.0955, 0.0977, 0.0980, 0.1007, 0.1063,
        0.1030],
        [0.0965, 0.0991, 0.1017, 0.1013, 0.0958, 0.0979, 0.0978, 0.1004, 0.1067,
        0.1029],
        [0.0965, 0.1000, 0.1018, 0.1011, 0.0956, 0.0982, 0.0980, 0.1005, 0.1055,
        0.1029],
        [0.0966, 0.0996, 0.1017, 0.1010, 0.0956, 0.0981, 0.0980, 0.1005, 0.1063,
        0.1026],

```

```
[0.0968, 0.0994, 0.1016, 0.1015, 0.0957, 0.0974, 0.0977, 0.1006, 0.1067,
 0.1024],
[0.0962, 0.0996, 0.1018, 0.1010, 0.0958, 0.0979, 0.0978, 0.1009, 0.1062,
 0.1027],
[0.0970, 0.0992, 0.1017, 0.1016, 0.0953, 0.0982, 0.0978, 0.1002, 0.1061,
 0.1029],
[0.0963, 0.0997, 0.1019, 0.1014, 0.0954, 0.0979, 0.0982, 0.1006, 0.1062,
 0.1025],
[0.0973, 0.0990, 0.1017, 0.1015, 0.0954, 0.0975, 0.0979, 0.1001, 0.1068,
 0.1028],
[0.0965, 0.0990, 0.1017, 0.1014, 0.0955, 0.0982, 0.0976, 0.1006, 0.1065,
 0.1030],
[0.0965, 0.0994, 0.1020, 0.1015, 0.0956, 0.0977, 0.0979, 0.1004, 0.1061,
 0.1030],
[0.0967, 0.0993, 0.1018, 0.1016, 0.0955, 0.0976, 0.0980, 0.1006, 0.1062,
 0.1027],
[0.0968, 0.0995, 0.1018, 0.1013, 0.0954, 0.0983, 0.0979, 0.1003, 0.1064,
 0.1025],
[0.0965, 0.0993, 0.1014, 0.1016, 0.0957, 0.0973, 0.0975, 0.1004, 0.1068,
 0.1033],
[0.0962, 0.0994, 0.1020, 0.1015, 0.0952, 0.0982, 0.0981, 0.1003, 0.1065,
 0.1026],
[0.0963, 0.0992, 0.1015, 0.1014, 0.0956, 0.0974, 0.0978, 0.1010, 0.1066,
 0.1032],
[0.0962, 0.0998, 0.1016, 0.1011, 0.0959, 0.0980, 0.0976, 0.1007, 0.1062,
 0.1029],
[0.0965, 0.0994, 0.1017, 0.1014, 0.0957, 0.0980, 0.0980, 0.1007, 0.1061,
 0.1027],
[0.0960, 0.0994, 0.1016, 0.1011, 0.0957, 0.0980, 0.0983, 0.1004, 0.1067,
 0.1027],
[0.0963, 0.0995, 0.1016, 0.1014, 0.0960, 0.0975, 0.0978, 0.1007, 0.1063,
 0.1029]], device='cuda:0', grad_fn=<SoftmaxBackward>)
```

The following cell will define the training and the testing loop.

I have already defined the training loop to give you guys some idea. Try changing the learning rates, EPOCHS, weight\_decay, and optimizers etc. to see how the learning process changes.

Your task is to define the test function which prints the test accuracy, given the model and test\_loader. Do remember that while testing we use a single image (batch of size 1).

First understand the train function properly. Then attempt to write the test function.

(HINT: you don't need optimizer (torch.optim) and loss function (nn.CrossEntropyLoss) for testing).

```
[6]: def train(model, train_loader, EPOCHS = 6, lossF = None):

    if lossF == None:
```

```

    lossF = nn.CrossEntropyLoss() #Cross entropy loss is popularly used for
    ↪classification tasks.

    ## Adam is a very popular choice of optimization algorithms.
    ## It is not very sensitive to hyperparameters and therefore it becomes a
    ↪natural choice in quick experiments.
    optim = torch.optim.Adam (model.parameters(), lr = 4e-4, weight_decay=1e-3)

    model.train() #Changes the model to train mode. All the require_grad are set
    ↪to true.

    for epoch in range(EPOCHS):
        correct = 0
        for batch_idx, (X_batch, y_batch) in enumerate(train_loader):

            #Move data to device
            var_X_batch = Variable(X_batch).to(device)
            var_y_batch = Variable(y_batch).to(device)
            # print(var_X_batch.size())

            ## Forward Pass!
            output = model(var_X_batch)

            #print(output)
            #print(output.size(), var_y_batch.size())

            ## Calculate the loss incurred
            loss = lossF (output, var_y_batch)

            ## BackProp: Computes all gradients.
            loss.backward()

            ## Gradient Descent Step (Adam)
            optim.step()
            optim.zero_grad() # This is important because PyTorch keeps on adding to
            ↪the original value of gradient.

            ## Gets the predictions. From probabilities (the digit with highest
            ↪probability is the prediction)
            predicted = torch.max(output.data, axis = 1).indices

            # print(predicted)

            ## Calculates the number of correct predictions in a batch
            correct += (predicted == var_y_batch).sum()

        if (batch_idx % 200) == 0:

```



```

        print('Epoch : {} [{} / {} ( {:.0f} % )] \t Loss: {:.6f} \t Accuracy: {:.3f} %'.
        ↪ format(
            epoch, batch_idx * len(X_batch), len(train_loader.dataset),
        ↪ 100. * batch_idx / len(train_loader), loss.item(), float(correct * 100) /
        ↪ float(BATCH_SIZE * (batch_idx + 1)))

def test(model, test_loader):
    """
    Change the value of correct to contain the number of correctly classified
    ↪ test examples out of the 10000 example in test_loader.
    """
    correct = 0
    # Put your code here:
    # __begin
    for test_imgs, test_labels in test_loader:
        # print(test_imgs.shape)
        test_imgs = Variable(test_imgs).float()
        test_imgs = Variable(test_imgs).to(device)
        test_labels = Variable(test_labels).to(device)
        output = model(test_imgs)
        predicted = torch.max(output, 1)[1]
        correct += (predicted == test_labels).sum()

    # __end
    print("Test accuracy: {:.3f} % ".format( float(correct * 100) /
    ↪ (len(test_loader))))

```

Running the following cell will start training of the "net".

Note: Any decent bugfree conv-net should easily reach an accuracy of atleast 96%. If you are unable to do so, try improving the network architecture and/or try different training rate or longer training (higher EPOCHS)

```

[7]: net = net.cuda()
      train (net, train_loader, EPOCHS = 20)

```

Epoch : 0 [0/60000 (0%)]	Loss: 2.302599	Accuracy: 6.250%
Epoch : 0 [6400/60000 (11%)]	Loss: 1.784655	Accuracy: 53.887%
Epoch : 0 [12800/60000 (21%)]	Loss: 1.744274	Accuracy: 62.851%
Epoch : 0 [19200/60000 (32%)]	Loss: 1.681714	Accuracy: 68.480%
Epoch : 0 [25600/60000 (43%)]	Loss: 1.576507	Accuracy: 73.490%
Epoch : 0 [32000/60000 (53%)]	Loss: 1.494225	Accuracy: 77.123%
Epoch : 0 [38400/60000 (64%)]	Loss: 1.506610	Accuracy: 79.874%
Epoch : 0 [44800/60000 (75%)]	Loss: 1.487906	Accuracy: 81.883%
Epoch : 0 [51200/60000 (85%)]	Loss: 1.527602	Accuracy: 83.403%
Epoch : 0 [57600/60000 (96%)]	Loss: 1.514520	Accuracy: 84.708%

Epoch : 1	[0/60000 (0%)]	Loss: 1.488366	Accuracy:100.000%
Epoch : 1	[6400/60000 (11%)]	Loss: 1.502373	Accuracy:96.160%
Epoch : 1	[12800/60000 (21%)]	Loss: 1.475746	Accuracy:95.776%
Epoch : 1	[19200/60000 (32%)]	Loss: 1.479310	Accuracy:95.726%
Epoch : 1	[25600/60000 (43%)]	Loss: 1.464445	Accuracy:95.876%
Epoch : 1	[32000/60000 (53%)]	Loss: 1.498445	Accuracy:95.970%
Epoch : 1	[38400/60000 (64%)]	Loss: 1.492334	Accuracy:96.074%
Epoch : 1	[44800/60000 (75%)]	Loss: 1.491384	Accuracy:96.117%
Epoch : 1	[51200/60000 (85%)]	Loss: 1.490910	Accuracy:96.110%
Epoch : 1	[57600/60000 (96%)]	Loss: 1.500375	Accuracy:96.167%
Epoch : 2	[0/60000 (0%)]	Loss: 1.471059	Accuracy:100.000%
Epoch : 2	[6400/60000 (11%)]	Loss: 1.462476	Accuracy:97.248%
Epoch : 2	[12800/60000 (21%)]	Loss: 1.476719	Accuracy:96.820%
Epoch : 2	[19200/60000 (32%)]	Loss: 1.479044	Accuracy:96.823%
Epoch : 2	[25600/60000 (43%)]	Loss: 1.482142	Accuracy:97.000%
Epoch : 2	[32000/60000 (53%)]	Loss: 1.493994	Accuracy:96.994%
Epoch : 2	[38400/60000 (64%)]	Loss: 1.521864	Accuracy:97.042%
Epoch : 2	[44800/60000 (75%)]	Loss: 1.476910	Accuracy:97.047%
Epoch : 2	[51200/60000 (85%)]	Loss: 1.492091	Accuracy:97.016%
Epoch : 2	[57600/60000 (96%)]	Loss: 1.505467	Accuracy:97.050%
Epoch : 3	[0/60000 (0%)]	Loss: 1.461615	Accuracy:100.000%
Epoch : 3	[6400/60000 (11%)]	Loss: 1.465881	Accuracy:97.404%
Epoch : 3	[12800/60000 (21%)]	Loss: 1.494033	Accuracy:97.195%
Epoch : 3	[19200/60000 (32%)]	Loss: 1.463103	Accuracy:97.270%
Epoch : 3	[25600/60000 (43%)]	Loss: 1.477686	Accuracy:97.413%
Epoch : 3	[32000/60000 (53%)]	Loss: 1.462061	Accuracy:97.418%
Epoch : 3	[38400/60000 (64%)]	Loss: 1.513970	Accuracy:97.440%
Epoch : 3	[44800/60000 (75%)]	Loss: 1.492633	Accuracy:97.455%
Epoch : 3	[51200/60000 (85%)]	Loss: 1.489715	Accuracy:97.422%
Epoch : 3	[57600/60000 (96%)]	Loss: 1.497450	Accuracy:97.399%
Epoch : 4	[0/60000 (0%)]	Loss: 1.461385	Accuracy:100.000%
Epoch : 4	[6400/60000 (11%)]	Loss: 1.490536	Accuracy:97.761%
Epoch : 4	[12800/60000 (21%)]	Loss: 1.529633	Accuracy:97.576%
Epoch : 4	[19200/60000 (32%)]	Loss: 1.462422	Accuracy:97.702%
Epoch : 4	[25600/60000 (43%)]	Loss: 1.467066	Accuracy:97.772%
Epoch : 4	[32000/60000 (53%)]	Loss: 1.472054	Accuracy:97.727%
Epoch : 4	[38400/60000 (64%)]	Loss: 1.491543	Accuracy:97.741%
Epoch : 4	[44800/60000 (75%)]	Loss: 1.493542	Accuracy:97.743%
Epoch : 4	[51200/60000 (85%)]	Loss: 1.492354	Accuracy:97.689%
Epoch : 4	[57600/60000 (96%)]	Loss: 1.492149	Accuracy:97.701%
Epoch : 5	[0/60000 (0%)]	Loss: 1.463364	Accuracy:100.000%
Epoch : 5	[6400/60000 (11%)]	Loss: 1.464313	Accuracy:98.010%
Epoch : 5	[12800/60000 (21%)]	Loss: 1.502812	Accuracy:97.857%
Epoch : 5	[19200/60000 (32%)]	Loss: 1.466466	Accuracy:97.837%
Epoch : 5	[25600/60000 (43%)]	Loss: 1.468061	Accuracy:97.885%
Epoch : 5	[32000/60000 (53%)]	Loss: 1.479162	Accuracy:97.890%
Epoch : 5	[38400/60000 (64%)]	Loss: 1.491776	Accuracy:97.926%
Epoch : 5	[44800/60000 (75%)]	Loss: 1.463630	Accuracy:97.952%

Epoch : 5	[51200/60000 (85%)]	Loss: 1.489880	Accuracy:97.908%
Epoch : 5	[57600/60000 (96%)]	Loss: 1.488033	Accuracy:97.928%
Epoch : 6	[0/60000 (0%)]	Loss: 1.462193	Accuracy:100.000%
Epoch : 6	[6400/60000 (11%)]	Loss: 1.466863	Accuracy:97.839%
Epoch : 6	[12800/60000 (21%)]	Loss: 1.483182	Accuracy:97.795%
Epoch : 6	[19200/60000 (32%)]	Loss: 1.473910	Accuracy:97.910%
Epoch : 6	[25600/60000 (43%)]	Loss: 1.464815	Accuracy:97.956%
Epoch : 6	[32000/60000 (53%)]	Loss: 1.487100	Accuracy:98.002%
Epoch : 6	[38400/60000 (64%)]	Loss: 1.491885	Accuracy:98.017%
Epoch : 6	[44800/60000 (75%)]	Loss: 1.466469	Accuracy:98.035%
Epoch : 6	[51200/60000 (85%)]	Loss: 1.491591	Accuracy:97.982%
Epoch : 6	[57600/60000 (96%)]	Loss: 1.483985	Accuracy:98.015%
Epoch : 7	[0/60000 (0%)]	Loss: 1.461212	Accuracy:100.000%
Epoch : 7	[6400/60000 (11%)]	Loss: 1.464385	Accuracy:97.963%
Epoch : 7	[12800/60000 (21%)]	Loss: 1.477983	Accuracy:98.021%
Epoch : 7	[19200/60000 (32%)]	Loss: 1.469435	Accuracy:98.102%
Epoch : 7	[25600/60000 (43%)]	Loss: 1.465861	Accuracy:98.186%
Epoch : 7	[32000/60000 (53%)]	Loss: 1.498821	Accuracy:98.186%
Epoch : 7	[38400/60000 (64%)]	Loss: 1.492798	Accuracy:98.189%
Epoch : 7	[44800/60000 (75%)]	Loss: 1.464990	Accuracy:98.178%
Epoch : 7	[51200/60000 (85%)]	Loss: 1.491070	Accuracy:98.113%
Epoch : 7	[57600/60000 (96%)]	Loss: 1.492181	Accuracy:98.140%
Epoch : 8	[0/60000 (0%)]	Loss: 1.461977	Accuracy:100.000%
Epoch : 8	[6400/60000 (11%)]	Loss: 1.462248	Accuracy:98.212%
Epoch : 8	[12800/60000 (21%)]	Loss: 1.489834	Accuracy:98.130%
Epoch : 8	[19200/60000 (32%)]	Loss: 1.465303	Accuracy:98.201%
Epoch : 8	[25600/60000 (43%)]	Loss: 1.462251	Accuracy:98.283%
Epoch : 8	[32000/60000 (53%)]	Loss: 1.490436	Accuracy:98.308%
Epoch : 8	[38400/60000 (64%)]	Loss: 1.491498	Accuracy:98.309%
Epoch : 8	[44800/60000 (75%)]	Loss: 1.461868	Accuracy:98.320%
Epoch : 8	[51200/60000 (85%)]	Loss: 1.491032	Accuracy:98.257%
Epoch : 8	[57600/60000 (96%)]	Loss: 1.473637	Accuracy:98.263%
Epoch : 9	[0/60000 (0%)]	Loss: 1.461248	Accuracy:100.000%
Epoch : 9	[6400/60000 (11%)]	Loss: 1.473389	Accuracy:98.321%
Epoch : 9	[12800/60000 (21%)]	Loss: 1.482455	Accuracy:98.262%
Epoch : 9	[19200/60000 (32%)]	Loss: 1.462209	Accuracy:98.305%
Epoch : 9	[25600/60000 (43%)]	Loss: 1.463193	Accuracy:98.389%
Epoch : 9	[32000/60000 (53%)]	Loss: 1.462231	Accuracy:98.367%
Epoch : 9	[38400/60000 (64%)]	Loss: 1.503498	Accuracy:98.361%
Epoch : 9	[44800/60000 (75%)]	Loss: 1.461265	Accuracy:98.378%
Epoch : 9	[51200/60000 (85%)]	Loss: 1.490982	Accuracy:98.317%
Epoch : 9	[57600/60000 (96%)]	Loss: 1.477924	Accuracy:98.327%
Epoch : 10	[0/60000 (0%)]	Loss: 1.461547	Accuracy:100.000%
Epoch : 10	[6400/60000 (11%)]	Loss: 1.479868	Accuracy:98.305%
Epoch : 10	[12800/60000 (21%)]	Loss: 1.498153	Accuracy:98.301%
Epoch : 10	[19200/60000 (32%)]	Loss: 1.464441	Accuracy:98.367%
Epoch : 10	[25600/60000 (43%)]	Loss: 1.467145	Accuracy:98.404%
Epoch : 10	[32000/60000 (53%)]	Loss: 1.484795	Accuracy:98.358%

Epoch : 10	[38400/60000 (64%)]	Loss: 1.490959	Accuracy:98.371%
Epoch : 10	[44800/60000 (75%)]	Loss: 1.463867	Accuracy:98.387%
Epoch : 10	[51200/60000 (85%)]	Loss: 1.492266	Accuracy:98.333%
Epoch : 10	[57600/60000 (96%)]	Loss: 1.467137	Accuracy:98.367%
Epoch : 11	[0/60000 (0%)]	Loss: 1.461564	Accuracy:100.000%
Epoch : 11	[6400/60000 (11%)]	Loss: 1.462870	Accuracy:98.274%
Epoch : 11	[12800/60000 (21%)]	Loss: 1.483328	Accuracy:98.325%
Epoch : 11	[19200/60000 (32%)]	Loss: 1.462031	Accuracy:98.425%
Epoch : 11	[25600/60000 (43%)]	Loss: 1.462825	Accuracy:98.486%
Epoch : 11	[32000/60000 (53%)]	Loss: 1.463807	Accuracy:98.483%
Epoch : 11	[38400/60000 (64%)]	Loss: 1.491063	Accuracy:98.486%
Epoch : 11	[44800/60000 (75%)]	Loss: 1.471956	Accuracy:98.499%
Epoch : 11	[51200/60000 (85%)]	Loss: 1.467975	Accuracy:98.444%
Epoch : 11	[57600/60000 (96%)]	Loss: 1.470811	Accuracy:98.461%
Epoch : 12	[0/60000 (0%)]	Loss: 1.462924	Accuracy:100.000%
Epoch : 12	[6400/60000 (11%)]	Loss: 1.462150	Accuracy:98.539%
Epoch : 12	[12800/60000 (21%)]	Loss: 1.485432	Accuracy:98.473%
Epoch : 12	[19200/60000 (32%)]	Loss: 1.466822	Accuracy:98.513%
Epoch : 12	[25600/60000 (43%)]	Loss: 1.475937	Accuracy:98.572%
Epoch : 12	[32000/60000 (53%)]	Loss: 1.463308	Accuracy:98.601%
Epoch : 12	[38400/60000 (64%)]	Loss: 1.490987	Accuracy:98.579%
Epoch : 12	[44800/60000 (75%)]	Loss: 1.483012	Accuracy:98.588%
Epoch : 12	[51200/60000 (85%)]	Loss: 1.488576	Accuracy:98.546%
Epoch : 12	[57600/60000 (96%)]	Loss: 1.498243	Accuracy:98.562%
Epoch : 13	[0/60000 (0%)]	Loss: 1.461428	Accuracy:100.000%
Epoch : 13	[6400/60000 (11%)]	Loss: 1.467439	Accuracy:98.445%
Epoch : 13	[12800/60000 (21%)]	Loss: 1.482636	Accuracy:98.340%
Epoch : 13	[19200/60000 (32%)]	Loss: 1.461794	Accuracy:98.419%
Epoch : 13	[25600/60000 (43%)]	Loss: 1.461940	Accuracy:98.537%
Epoch : 13	[32000/60000 (53%)]	Loss: 1.476435	Accuracy:98.536%
Epoch : 13	[38400/60000 (64%)]	Loss: 1.490933	Accuracy:98.564%
Epoch : 13	[44800/60000 (75%)]	Loss: 1.483143	Accuracy:98.550%
Epoch : 13	[51200/60000 (85%)]	Loss: 1.470785	Accuracy:98.546%
Epoch : 13	[57600/60000 (96%)]	Loss: 1.485798	Accuracy:98.569%
Epoch : 14	[0/60000 (0%)]	Loss: 1.461482	Accuracy:100.000%
Epoch : 14	[6400/60000 (11%)]	Loss: 1.464975	Accuracy:98.492%
Epoch : 14	[12800/60000 (21%)]	Loss: 1.473536	Accuracy:98.543%
Epoch : 14	[19200/60000 (32%)]	Loss: 1.462132	Accuracy:98.560%
Epoch : 14	[25600/60000 (43%)]	Loss: 1.464996	Accuracy:98.607%
Epoch : 14	[32000/60000 (53%)]	Loss: 1.463661	Accuracy:98.614%
Epoch : 14	[38400/60000 (64%)]	Loss: 1.491048	Accuracy:98.631%
Epoch : 14	[44800/60000 (75%)]	Loss: 1.464110	Accuracy:98.601%
Epoch : 14	[51200/60000 (85%)]	Loss: 1.486419	Accuracy:98.563%
Epoch : 14	[57600/60000 (96%)]	Loss: 1.481739	Accuracy:98.581%
Epoch : 15	[0/60000 (0%)]	Loss: 1.462240	Accuracy:100.000%
Epoch : 15	[6400/60000 (11%)]	Loss: 1.472030	Accuracy:98.756%
Epoch : 15	[12800/60000 (21%)]	Loss: 1.480627	Accuracy:98.667%
Epoch : 15	[19200/60000 (32%)]	Loss: 1.464678	Accuracy:98.705%

Epoch : 15	[25600/60000 (43%)]	Loss: 1.464290	Accuracy:98.755%
Epoch : 15	[32000/60000 (53%)]	Loss: 1.464509	Accuracy:98.748%
Epoch : 15	[38400/60000 (64%)]	Loss: 1.491058	Accuracy:98.733%
Epoch : 15	[44800/60000 (75%)]	Loss: 1.465672	Accuracy:98.711%
Epoch : 15	[51200/60000 (85%)]	Loss: 1.479563	Accuracy:98.657%
Epoch : 15	[57600/60000 (96%)]	Loss: 1.466746	Accuracy:98.659%
Epoch : 16	[0/60000 (0%)]	Loss: 1.461531	Accuracy:100.000%
Epoch : 16	[6400/60000 (11%)]	Loss: 1.468590	Accuracy:98.616%
Epoch : 16	[12800/60000 (21%)]	Loss: 1.497316	Accuracy:98.550%
Epoch : 16	[19200/60000 (32%)]	Loss: 1.466259	Accuracy:98.648%
Epoch : 16	[25600/60000 (43%)]	Loss: 1.461793	Accuracy:98.713%
Epoch : 16	[32000/60000 (53%)]	Loss: 1.466365	Accuracy:98.708%
Epoch : 16	[38400/60000 (64%)]	Loss: 1.491099	Accuracy:98.686%
Epoch : 16	[44800/60000 (75%)]	Loss: 1.471007	Accuracy:98.684%
Epoch : 16	[51200/60000 (85%)]	Loss: 1.490944	Accuracy:98.645%
Epoch : 16	[57600/60000 (96%)]	Loss: 1.488891	Accuracy:98.662%
Epoch : 17	[0/60000 (0%)]	Loss: 1.466717	Accuracy:100.000%
Epoch : 17	[6400/60000 (11%)]	Loss: 1.472402	Accuracy:98.741%
Epoch : 17	[12800/60000 (21%)]	Loss: 1.463979	Accuracy:98.488%
Epoch : 17	[19200/60000 (32%)]	Loss: 1.463227	Accuracy:98.601%
Epoch : 17	[25600/60000 (43%)]	Loss: 1.463009	Accuracy:98.674%
Epoch : 17	[32000/60000 (53%)]	Loss: 1.463393	Accuracy:98.689%
Epoch : 17	[38400/60000 (64%)]	Loss: 1.498358	Accuracy:98.678%
Epoch : 17	[44800/60000 (75%)]	Loss: 1.464586	Accuracy:98.695%
Epoch : 17	[51200/60000 (85%)]	Loss: 1.488403	Accuracy:98.657%
Epoch : 17	[57600/60000 (96%)]	Loss: 1.464547	Accuracy:98.681%
Epoch : 18	[0/60000 (0%)]	Loss: 1.461233	Accuracy:100.000%
Epoch : 18	[6400/60000 (11%)]	Loss: 1.462193	Accuracy:98.601%
Epoch : 18	[12800/60000 (21%)]	Loss: 1.472664	Accuracy:98.527%
Epoch : 18	[19200/60000 (32%)]	Loss: 1.465309	Accuracy:98.612%
Epoch : 18	[25600/60000 (43%)]	Loss: 1.462420	Accuracy:98.685%
Epoch : 18	[32000/60000 (53%)]	Loss: 1.462830	Accuracy:98.698%
Epoch : 18	[38400/60000 (64%)]	Loss: 1.490883	Accuracy:98.712%
Epoch : 18	[44800/60000 (75%)]	Loss: 1.469108	Accuracy:98.706%
Epoch : 18	[51200/60000 (85%)]	Loss: 1.472325	Accuracy:98.673%
Epoch : 18	[57600/60000 (96%)]	Loss: 1.466831	Accuracy:98.688%
Epoch : 19	[0/60000 (0%)]	Loss: 1.461330	Accuracy:100.000%
Epoch : 19	[6400/60000 (11%)]	Loss: 1.461505	Accuracy:98.865%
Epoch : 19	[12800/60000 (21%)]	Loss: 1.474587	Accuracy:98.776%
Epoch : 19	[19200/60000 (32%)]	Loss: 1.463233	Accuracy:98.788%
Epoch : 19	[25600/60000 (43%)]	Loss: 1.464973	Accuracy:98.869%
Epoch : 19	[32000/60000 (53%)]	Loss: 1.468030	Accuracy:98.842%
Epoch : 19	[38400/60000 (64%)]	Loss: 1.491192	Accuracy:98.834%
Epoch : 19	[44800/60000 (75%)]	Loss: 1.491296	Accuracy:98.827%
Epoch : 19	[51200/60000 (85%)]	Loss: 1.487845	Accuracy:98.761%
Epoch : 19	[57600/60000 (96%)]	Loss: 1.465441	Accuracy:98.782%

Run the test function that you defined above!

Note: You should easily reach 95%.

```
[8]: test(net, test_loader)
```

Test accuracy:98.570%

```
[29]: # My second attempt, has only 2 conv layers and performs nearly as well
```

```
class Model2(torch.nn.Module):
    def __init__(self):
        super(Model2, self).__init__()
        self.conv_1 = torch.nn.Conv2d(1, 32, 3, stride=1, padding=1)
        self.conv_2 = torch.nn.Conv2d(32, 64, 3, stride=1, padding=1)
        self.max_pool2d = torch.nn.MaxPool2d(kernel_size=2, stride=2)
        self.linear_1 = torch.nn.Linear(7 * 7 * 64, 128)
        self.linear_2 = torch.nn.Linear(128, 10)
        self.relu = torch.nn.ReLU()

    def forward(self, x):
        x = self.conv_1(x)
        x = self.relu(x)
        x = self.max_pool2d(x)

        x = self.conv_2(x)
        x = self.relu(x)
        x = self.max_pool2d(x)

        x = x.reshape(x.size(0), -1)

        x = self.linear_1(x)
        x = self.relu(x)

        x = self.linear_2(x)

        pred = F.softmax(x, dim=1)

        return pred
```

```
[30]: net2 = Model2()
net2.cuda()
net2.to(device)
print(net2)
tld = iter(train_loader)
im = next(tld)[0].to(device)
print('Size of the output tensor:', net2.forward(im).size())
print(net2.forward(im))
```

Model2(

```

(conv_1): Conv2d(1, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(conv_2): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(max_pool2d): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(linear_1): Linear(in_features=3136, out_features=128, bias=True)
(linear_2): Linear(in_features=128, out_features=10, bias=True)
(relu): ReLU()
)
Size of the output tensor: torch.Size([32, 10])
tensor([[0.1015, 0.1008, 0.0984, 0.1059, 0.0940, 0.0907, 0.1061, 0.1047, 0.0968,
0.1009],
[0.1011, 0.1032, 0.0976, 0.1044, 0.0927, 0.0920, 0.1068, 0.1061, 0.0957,
0.1005],
[0.1014, 0.1045, 0.0941, 0.1026, 0.0948, 0.0934, 0.1050, 0.1046, 0.0967,
0.1030],
[0.0984, 0.1037, 0.0992, 0.1032, 0.0936, 0.0926, 0.1045, 0.1049, 0.0971,
0.1028],
[0.0985, 0.1020, 0.0979, 0.1049, 0.0943, 0.0920, 0.1073, 0.1050, 0.0967,
0.1014],
[0.0988, 0.0998, 0.0983, 0.1056, 0.0948, 0.0916, 0.1064, 0.1061, 0.0975,
0.1011],
[0.1014, 0.0988, 0.0990, 0.1045, 0.0926, 0.0924, 0.1066, 0.1064, 0.0960,
0.1023],
[0.1006, 0.1002, 0.0983, 0.1053, 0.0938, 0.0916, 0.1063, 0.1072, 0.0962,
0.1006],
[0.1014, 0.1011, 0.0993, 0.1041, 0.0918, 0.0917, 0.1064, 0.1042, 0.0969,
0.1031],
[0.0993, 0.1024, 0.0984, 0.1040, 0.0943, 0.0935, 0.1052, 0.1046, 0.0965,
0.1018],
[0.1011, 0.0995, 0.0986, 0.1045, 0.0936, 0.0928, 0.1062, 0.1071, 0.0952,
0.1014],
[0.0988, 0.1035, 0.0986, 0.1035, 0.0943, 0.0929, 0.1045, 0.1057, 0.0961,
0.1023],
[0.1009, 0.1014, 0.0963, 0.1053, 0.0952, 0.0919, 0.1056, 0.1066, 0.0951,
0.1017],
[0.1000, 0.0999, 0.0982, 0.1050, 0.0934, 0.0949, 0.1060, 0.1072, 0.0940,
0.1014],
[0.1016, 0.1013, 0.0991, 0.1037, 0.0921, 0.0913, 0.1060, 0.1043, 0.0972,
0.1034],
[0.0999, 0.1021, 0.0979, 0.1041, 0.0935, 0.0923, 0.1073, 0.1049, 0.0967,
0.1014],
[0.1008, 0.0992, 0.0981, 0.1043, 0.0940, 0.0930, 0.1071, 0.1054, 0.0965,
0.1016],
[0.0999, 0.1013, 0.0996, 0.1049, 0.0939, 0.0919, 0.1071, 0.1049, 0.0958,
0.1007],
[0.1005, 0.0983, 0.0990, 0.1057, 0.0943, 0.0933, 0.1054, 0.1062, 0.0952,
0.1022],
[0.0996, 0.1013, 0.0991, 0.1053, 0.0934, 0.0922, 0.1059, 0.1051, 0.0962,

```

```

0.1019],
[0.0971, 0.1032, 0.0963, 0.1034, 0.0956, 0.0947, 0.1052, 0.1063, 0.0954,
0.1027],
[0.1005, 0.1049, 0.0980, 0.1042, 0.0934, 0.0916, 0.1053, 0.1051, 0.0954,
0.1016],
[0.0987, 0.1000, 0.0988, 0.1058, 0.0945, 0.0923, 0.1063, 0.1049, 0.0967,
0.1020],
[0.0989, 0.1033, 0.0990, 0.1032, 0.0938, 0.0926, 0.1048, 0.1051, 0.0967,
0.1028],
[0.1002, 0.1011, 0.0987, 0.1038, 0.0935, 0.0928, 0.1052, 0.1060, 0.0962,
0.1024],
[0.0994, 0.1012, 0.1002, 0.1045, 0.0932, 0.0913, 0.1063, 0.1066, 0.0961,
0.1013],
[0.0994, 0.1011, 0.0971, 0.1056, 0.0945, 0.0924, 0.1063, 0.1047, 0.0965,
0.1023],
[0.1005, 0.1018, 0.0981, 0.1047, 0.0939, 0.0907, 0.1062, 0.1074, 0.0957,
0.1009],
[0.1004, 0.1021, 0.0980, 0.1042, 0.0936, 0.0911, 0.1076, 0.1065, 0.0953,
0.1013],
[0.0991, 0.1024, 0.0970, 0.1043, 0.0948, 0.0925, 0.1063, 0.1050, 0.0964,
0.1023],
[0.0995, 0.1000, 0.0968, 0.1054, 0.0932, 0.0938, 0.1058, 0.1064, 0.0965,
0.1027],
[0.0989, 0.1014, 0.0994, 0.1039, 0.0940, 0.0922, 0.1069, 0.1051, 0.0971,
0.1011]], device='cuda:0', grad_fn=<SoftmaxBackward>)

```

```
[31]: train (net2, train_loader, EPOCHS = 20)
```

```

Epoch : 0 [0/60000 (0%)]      Loss: 2.302189   Accuracy:0.000%
Epoch : 0 [6400/60000 (11%)]  Loss: 1.818810   Accuracy:55.442%
Epoch : 0 [12800/60000 (21%)] Loss: 1.626982   Accuracy:61.440%
Epoch : 0 [19200/60000 (32%)] Loss: 1.670997   Accuracy:67.549%
Epoch : 0 [25600/60000 (43%)] Loss: 1.559201   Accuracy:72.374%
Epoch : 0 [32000/60000 (53%)] Loss: 1.578661   Accuracy:76.580%
Epoch : 0 [38400/60000 (64%)] Loss: 1.516968   Accuracy:79.650%
Epoch : 0 [44800/60000 (75%)] Loss: 1.502508   Accuracy:81.794%
Epoch : 0 [51200/60000 (85%)] Loss: 1.515632   Accuracy:83.421%
Epoch : 0 [57600/60000 (96%)] Loss: 1.511342   Accuracy:84.793%
Epoch : 1 [0/60000 (0%)]      Loss: 1.495278   Accuracy:96.875%
Epoch : 1 [6400/60000 (11%)]  Loss: 1.493335   Accuracy:96.580%
Epoch : 1 [12800/60000 (21%)] Loss: 1.489852   Accuracy:96.080%
Epoch : 1 [19200/60000 (32%)] Loss: 1.482573   Accuracy:96.152%
Epoch : 1 [25600/60000 (43%)] Loss: 1.462774   Accuracy:96.391%
Epoch : 1 [32000/60000 (53%)] Loss: 1.529930   Accuracy:96.441%
Epoch : 1 [38400/60000 (64%)] Loss: 1.513393   Accuracy:96.498%
Epoch : 1 [44800/60000 (75%)] Loss: 1.495590   Accuracy:96.525%
Epoch : 1 [51200/60000 (85%)] Loss: 1.488971   Accuracy:96.514%
Epoch : 1 [57600/60000 (96%)] Loss: 1.499197   Accuracy:96.590%

```



Epoch : 2	[0/60000 (0%)]	Loss: 1.465693	Accuracy:100.000%
Epoch : 2	[6400/60000 (11%)]	Loss: 1.490964	Accuracy:97.233%
Epoch : 2	[12800/60000 (21%)]	Loss: 1.472825	Accuracy:96.937%
Epoch : 2	[19200/60000 (32%)]	Loss: 1.471897	Accuracy:97.010%
Epoch : 2	[25600/60000 (43%)]	Loss: 1.502457	Accuracy:97.195%
Epoch : 2	[32000/60000 (53%)]	Loss: 1.523909	Accuracy:97.203%
Epoch : 2	[38400/60000 (64%)]	Loss: 1.509187	Accuracy:97.198%
Epoch : 2	[44800/60000 (75%)]	Loss: 1.489578	Accuracy:97.198%
Epoch : 2	[51200/60000 (85%)]	Loss: 1.490946	Accuracy:97.172%
Epoch : 2	[57600/60000 (96%)]	Loss: 1.490612	Accuracy:97.208%
Epoch : 3	[0/60000 (0%)]	Loss: 1.464954	Accuracy:100.000%
Epoch : 3	[6400/60000 (11%)]	Loss: 1.493724	Accuracy:97.404%
Epoch : 3	[12800/60000 (21%)]	Loss: 1.480821	Accuracy:97.296%
Epoch : 3	[19200/60000 (32%)]	Loss: 1.470373	Accuracy:97.333%
Epoch : 3	[25600/60000 (43%)]	Loss: 1.504097	Accuracy:97.538%
Epoch : 3	[32000/60000 (53%)]	Loss: 1.534427	Accuracy:97.549%
Epoch : 3	[38400/60000 (64%)]	Loss: 1.497270	Accuracy:97.552%
Epoch : 3	[44800/60000 (75%)]	Loss: 1.466729	Accuracy:97.535%
Epoch : 3	[51200/60000 (85%)]	Loss: 1.497637	Accuracy:97.470%
Epoch : 3	[57600/60000 (96%)]	Loss: 1.494090	Accuracy:97.500%
Epoch : 4	[0/60000 (0%)]	Loss: 1.463144	Accuracy:100.000%
Epoch : 4	[6400/60000 (11%)]	Loss: 1.496101	Accuracy:97.715%
Epoch : 4	[12800/60000 (21%)]	Loss: 1.488259	Accuracy:97.600%
Epoch : 4	[19200/60000 (32%)]	Loss: 1.474220	Accuracy:97.733%
Epoch : 4	[25600/60000 (43%)]	Loss: 1.487023	Accuracy:97.866%
Epoch : 4	[32000/60000 (53%)]	Loss: 1.529162	Accuracy:97.833%
Epoch : 4	[38400/60000 (64%)]	Loss: 1.495515	Accuracy:97.835%
Epoch : 4	[44800/60000 (75%)]	Loss: 1.473825	Accuracy:97.834%
Epoch : 4	[51200/60000 (85%)]	Loss: 1.498930	Accuracy:97.746%
Epoch : 4	[57600/60000 (96%)]	Loss: 1.482957	Accuracy:97.753%
Epoch : 5	[0/60000 (0%)]	Loss: 1.464787	Accuracy:100.000%
Epoch : 5	[6400/60000 (11%)]	Loss: 1.488373	Accuracy:97.823%
Epoch : 5	[12800/60000 (21%)]	Loss: 1.477630	Accuracy:97.693%
Epoch : 5	[19200/60000 (32%)]	Loss: 1.480989	Accuracy:97.821%
Epoch : 5	[25600/60000 (43%)]	Loss: 1.495892	Accuracy:97.960%
Epoch : 5	[32000/60000 (53%)]	Loss: 1.526620	Accuracy:97.958%
Epoch : 5	[38400/60000 (64%)]	Loss: 1.492473	Accuracy:97.963%
Epoch : 5	[44800/60000 (75%)]	Loss: 1.470722	Accuracy:97.957%
Epoch : 5	[51200/60000 (85%)]	Loss: 1.506606	Accuracy:97.904%
Epoch : 5	[57600/60000 (96%)]	Loss: 1.474086	Accuracy:97.892%
Epoch : 6	[0/60000 (0%)]	Loss: 1.464306	Accuracy:100.000%
Epoch : 6	[6400/60000 (11%)]	Loss: 1.490718	Accuracy:98.010%
Epoch : 6	[12800/60000 (21%)]	Loss: 1.477198	Accuracy:97.911%
Epoch : 6	[19200/60000 (32%)]	Loss: 1.480612	Accuracy:98.003%
Epoch : 6	[25600/60000 (43%)]	Loss: 1.483198	Accuracy:98.159%
Epoch : 6	[32000/60000 (53%)]	Loss: 1.506664	Accuracy:98.149%
Epoch : 6	[38400/60000 (64%)]	Loss: 1.491817	Accuracy:98.124%
Epoch : 6	[44800/60000 (75%)]	Loss: 1.469730	Accuracy:98.111%

Epoch : 6	[51200/60000 (85%)]	Loss: 1.503405	Accuracy:98.056%
Epoch : 6	[57600/60000 (96%)]	Loss: 1.481660	Accuracy:98.039%
Epoch : 7	[0/60000 (0%)]	Loss: 1.464681	Accuracy:100.000%
Epoch : 7	[6400/60000 (11%)]	Loss: 1.501098	Accuracy:98.057%
Epoch : 7	[12800/60000 (21%)]	Loss: 1.480210	Accuracy:98.013%
Epoch : 7	[19200/60000 (32%)]	Loss: 1.467210	Accuracy:98.066%
Epoch : 7	[25600/60000 (43%)]	Loss: 1.482053	Accuracy:98.194%
Epoch : 7	[32000/60000 (53%)]	Loss: 1.505728	Accuracy:98.199%
Epoch : 7	[38400/60000 (64%)]	Loss: 1.491651	Accuracy:98.207%
Epoch : 7	[44800/60000 (75%)]	Loss: 1.481989	Accuracy:98.207%
Epoch : 7	[51200/60000 (85%)]	Loss: 1.505931	Accuracy:98.173%
Epoch : 7	[57600/60000 (96%)]	Loss: 1.466335	Accuracy:98.142%
Epoch : 8	[0/60000 (0%)]	Loss: 1.464226	Accuracy:100.000%
Epoch : 8	[6400/60000 (11%)]	Loss: 1.490969	Accuracy:98.041%
Epoch : 8	[12800/60000 (21%)]	Loss: 1.486611	Accuracy:97.982%
Epoch : 8	[19200/60000 (32%)]	Loss: 1.469254	Accuracy:98.107%
Epoch : 8	[25600/60000 (43%)]	Loss: 1.485295	Accuracy:98.237%
Epoch : 8	[32000/60000 (53%)]	Loss: 1.506814	Accuracy:98.239%
Epoch : 8	[38400/60000 (64%)]	Loss: 1.491482	Accuracy:98.238%
Epoch : 8	[44800/60000 (75%)]	Loss: 1.480387	Accuracy:98.233%
Epoch : 8	[51200/60000 (85%)]	Loss: 1.510746	Accuracy:98.202%
Epoch : 8	[57600/60000 (96%)]	Loss: 1.470175	Accuracy:98.195%
Epoch : 9	[0/60000 (0%)]	Loss: 1.463546	Accuracy:100.000%
Epoch : 9	[6400/60000 (11%)]	Loss: 1.494410	Accuracy:98.181%
Epoch : 9	[12800/60000 (21%)]	Loss: 1.469811	Accuracy:98.067%
Epoch : 9	[19200/60000 (32%)]	Loss: 1.465618	Accuracy:98.206%
Epoch : 9	[25600/60000 (43%)]	Loss: 1.476494	Accuracy:98.291%
Epoch : 9	[32000/60000 (53%)]	Loss: 1.520020	Accuracy:98.305%
Epoch : 9	[38400/60000 (64%)]	Loss: 1.491564	Accuracy:98.319%
Epoch : 9	[44800/60000 (75%)]	Loss: 1.467684	Accuracy:98.305%
Epoch : 9	[51200/60000 (85%)]	Loss: 1.507861	Accuracy:98.278%
Epoch : 9	[57600/60000 (96%)]	Loss: 1.473866	Accuracy:98.274%
Epoch : 10	[0/60000 (0%)]	Loss: 1.463535	Accuracy:100.000%
Epoch : 10	[6400/60000 (11%)]	Loss: 1.506626	Accuracy:98.212%
Epoch : 10	[12800/60000 (21%)]	Loss: 1.474996	Accuracy:98.044%
Epoch : 10	[19200/60000 (32%)]	Loss: 1.468909	Accuracy:98.180%
Epoch : 10	[25600/60000 (43%)]	Loss: 1.479346	Accuracy:98.322%
Epoch : 10	[32000/60000 (53%)]	Loss: 1.515264	Accuracy:98.342%
Epoch : 10	[38400/60000 (64%)]	Loss: 1.491170	Accuracy:98.361%
Epoch : 10	[44800/60000 (75%)]	Loss: 1.466704	Accuracy:98.356%
Epoch : 10	[51200/60000 (85%)]	Loss: 1.501485	Accuracy:98.329%
Epoch : 10	[57600/60000 (96%)]	Loss: 1.480608	Accuracy:98.317%
Epoch : 11	[0/60000 (0%)]	Loss: 1.463524	Accuracy:100.000%
Epoch : 11	[6400/60000 (11%)]	Loss: 1.496235	Accuracy:98.197%
Epoch : 11	[12800/60000 (21%)]	Loss: 1.474311	Accuracy:98.044%
Epoch : 11	[19200/60000 (32%)]	Loss: 1.470642	Accuracy:98.243%
Epoch : 11	[25600/60000 (43%)]	Loss: 1.477164	Accuracy:98.346%
Epoch : 11	[32000/60000 (53%)]	Loss: 1.512678	Accuracy:98.370%

Epoch : 11	[38400/60000 (64%)]	Loss: 1.491341	Accuracy:98.410%
Epoch : 11	[44800/60000 (75%)]	Loss: 1.472550	Accuracy:98.394%
Epoch : 11	[51200/60000 (85%)]	Loss: 1.498510	Accuracy:98.380%
Epoch : 11	[57600/60000 (96%)]	Loss: 1.477948	Accuracy:98.367%
Epoch : 12	[0/60000 (0%)]	Loss: 1.464005	Accuracy:100.000%
Epoch : 12	[6400/60000 (11%)]	Loss: 1.506589	Accuracy:98.368%
Epoch : 12	[12800/60000 (21%)]	Loss: 1.478780	Accuracy:98.215%
Epoch : 12	[19200/60000 (32%)]	Loss: 1.467215	Accuracy:98.362%
Epoch : 12	[25600/60000 (43%)]	Loss: 1.476988	Accuracy:98.478%
Epoch : 12	[32000/60000 (53%)]	Loss: 1.515843	Accuracy:98.498%
Epoch : 12	[38400/60000 (64%)]	Loss: 1.491370	Accuracy:98.504%
Epoch : 12	[44800/60000 (75%)]	Loss: 1.472873	Accuracy:98.470%
Epoch : 12	[51200/60000 (85%)]	Loss: 1.500720	Accuracy:98.433%
Epoch : 12	[57600/60000 (96%)]	Loss: 1.471876	Accuracy:98.428%
Epoch : 13	[0/60000 (0%)]	Loss: 1.463105	Accuracy:100.000%
Epoch : 13	[6400/60000 (11%)]	Loss: 1.503983	Accuracy:98.336%
Epoch : 13	[12800/60000 (21%)]	Loss: 1.477224	Accuracy:98.270%
Epoch : 13	[19200/60000 (32%)]	Loss: 1.466545	Accuracy:98.352%
Epoch : 13	[25600/60000 (43%)]	Loss: 1.474158	Accuracy:98.471%
Epoch : 13	[32000/60000 (53%)]	Loss: 1.512237	Accuracy:98.498%
Epoch : 13	[38400/60000 (64%)]	Loss: 1.491049	Accuracy:98.509%
Epoch : 13	[44800/60000 (75%)]	Loss: 1.467018	Accuracy:98.479%
Epoch : 13	[51200/60000 (85%)]	Loss: 1.494388	Accuracy:98.454%
Epoch : 13	[57600/60000 (96%)]	Loss: 1.471699	Accuracy:98.445%
Epoch : 14	[0/60000 (0%)]	Loss: 1.462810	Accuracy:100.000%
Epoch : 14	[6400/60000 (11%)]	Loss: 1.499626	Accuracy:98.461%
Epoch : 14	[12800/60000 (21%)]	Loss: 1.482630	Accuracy:98.301%
Epoch : 14	[19200/60000 (32%)]	Loss: 1.466289	Accuracy:98.425%
Epoch : 14	[25600/60000 (43%)]	Loss: 1.475875	Accuracy:98.525%
Epoch : 14	[32000/60000 (53%)]	Loss: 1.508329	Accuracy:98.555%
Epoch : 14	[38400/60000 (64%)]	Loss: 1.491045	Accuracy:98.530%
Epoch : 14	[44800/60000 (75%)]	Loss: 1.469766	Accuracy:98.512%
Epoch : 14	[51200/60000 (85%)]	Loss: 1.507986	Accuracy:98.481%
Epoch : 14	[57600/60000 (96%)]	Loss: 1.474667	Accuracy:98.456%
Epoch : 15	[0/60000 (0%)]	Loss: 1.463211	Accuracy:100.000%
Epoch : 15	[6400/60000 (11%)]	Loss: 1.497269	Accuracy:98.585%
Epoch : 15	[12800/60000 (21%)]	Loss: 1.481877	Accuracy:98.457%
Epoch : 15	[19200/60000 (32%)]	Loss: 1.465198	Accuracy:98.523%
Epoch : 15	[25600/60000 (43%)]	Loss: 1.476220	Accuracy:98.607%
Epoch : 15	[32000/60000 (53%)]	Loss: 1.502374	Accuracy:98.633%
Epoch : 15	[38400/60000 (64%)]	Loss: 1.491117	Accuracy:98.621%
Epoch : 15	[44800/60000 (75%)]	Loss: 1.472667	Accuracy:98.588%
Epoch : 15	[51200/60000 (85%)]	Loss: 1.507521	Accuracy:98.556%
Epoch : 15	[57600/60000 (96%)]	Loss: 1.470615	Accuracy:98.546%
Epoch : 16	[0/60000 (0%)]	Loss: 1.462290	Accuracy:100.000%
Epoch : 16	[6400/60000 (11%)]	Loss: 1.496732	Accuracy:98.601%
Epoch : 16	[12800/60000 (21%)]	Loss: 1.487032	Accuracy:98.449%
Epoch : 16	[19200/60000 (32%)]	Loss: 1.465802	Accuracy:98.570%

Epoch : 16	[25600/60000 (43%)]	Loss: 1.476657	Accuracy:98.654%
Epoch : 16	[32000/60000 (53%)]	Loss: 1.502920	Accuracy:98.658%
Epoch : 16	[38400/60000 (64%)]	Loss: 1.491240	Accuracy:98.644%
Epoch : 16	[44800/60000 (75%)]	Loss: 1.465179	Accuracy:98.619%
Epoch : 16	[51200/60000 (85%)]	Loss: 1.483269	Accuracy:98.599%
Epoch : 16	[57600/60000 (96%)]	Loss: 1.471581	Accuracy:98.579%
Epoch : 17	[0/60000 (0%)]	Loss: 1.463269	Accuracy:100.000%
Epoch : 17	[6400/60000 (11%)]	Loss: 1.497206	Accuracy:98.523%
Epoch : 17	[12800/60000 (21%)]	Loss: 1.481294	Accuracy:98.465%
Epoch : 17	[19200/60000 (32%)]	Loss: 1.465512	Accuracy:98.554%
Epoch : 17	[25600/60000 (43%)]	Loss: 1.471454	Accuracy:98.662%
Epoch : 17	[32000/60000 (53%)]	Loss: 1.502186	Accuracy:98.654%
Epoch : 17	[38400/60000 (64%)]	Loss: 1.491244	Accuracy:98.652%
Epoch : 17	[44800/60000 (75%)]	Loss: 1.472111	Accuracy:98.626%
Epoch : 17	[51200/60000 (85%)]	Loss: 1.505261	Accuracy:98.612%
Epoch : 17	[57600/60000 (96%)]	Loss: 1.474675	Accuracy:98.600%
Epoch : 18	[0/60000 (0%)]	Loss: 1.462605	Accuracy:100.000%
Epoch : 18	[6400/60000 (11%)]	Loss: 1.494833	Accuracy:98.601%
Epoch : 18	[12800/60000 (21%)]	Loss: 1.493096	Accuracy:98.519%
Epoch : 18	[19200/60000 (32%)]	Loss: 1.466309	Accuracy:98.596%
Epoch : 18	[25600/60000 (43%)]	Loss: 1.475530	Accuracy:98.658%
Epoch : 18	[32000/60000 (53%)]	Loss: 1.504336	Accuracy:98.658%
Epoch : 18	[38400/60000 (64%)]	Loss: 1.491131	Accuracy:98.668%
Epoch : 18	[44800/60000 (75%)]	Loss: 1.469511	Accuracy:98.646%
Epoch : 18	[51200/60000 (85%)]	Loss: 1.502256	Accuracy:98.628%
Epoch : 18	[57600/60000 (96%)]	Loss: 1.470402	Accuracy:98.603%
Epoch : 19	[0/60000 (0%)]	Loss: 1.463854	Accuracy:100.000%
Epoch : 19	[6400/60000 (11%)]	Loss: 1.501996	Accuracy:98.601%
Epoch : 19	[12800/60000 (21%)]	Loss: 1.486060	Accuracy:98.504%
Epoch : 19	[19200/60000 (32%)]	Loss: 1.465211	Accuracy:98.622%
Epoch : 19	[25600/60000 (43%)]	Loss: 1.474832	Accuracy:98.701%
Epoch : 19	[32000/60000 (53%)]	Loss: 1.501074	Accuracy:98.711%
Epoch : 19	[38400/60000 (64%)]	Loss: 1.491320	Accuracy:98.696%
Epoch : 19	[44800/60000 (75%)]	Loss: 1.477404	Accuracy:98.642%
Epoch : 19	[51200/60000 (85%)]	Loss: 1.504792	Accuracy:98.608%
Epoch : 19	[57600/60000 (96%)]	Loss: 1.476464	Accuracy:98.607%

```
[33]: test(net2, test_loader)
```

Test accuracy:98.010%

## (Bonus) TASK 2

Before getting started read [this article](#) on Deep Dream.

Description of the task: Now we have trained the network to classify all the digits from 0 to 10. But is there a way to visualize what the network has learnt?

Well how about we learn (by backprop + Gradient Descent) the input image after fixing the expected output and all the network parameters. In other words say we want to visualize what 7

looks like to the network. Then we will do the following: 1. Initialize image tensor as `im` (maybe with all zeros). 2. Pass the image through the network. 3. Compute the loss with output of the net and expected output of 7. 4. Run backpropagation upto the image tensor. 5. Update the image using the update rule. (Specified by `torch.optim` object)

Remember that while creating the optimization object (`torch.optim` object) you have to pass to it a list of parameters to be optimized. Which in this case won't be `model.parameters()` but rather `[im, ]`.

Note: We are not using training or test data anywhere here!

```
[34]: def train_im(model, train_loader, digit = 7 ,iters = 1000, lossF = None):
      """
      Train the input image to match the <digit>. Run <iters> iterations of
      ↪Gradient Descent.
      """
      im = torch.zeros_like(train_data[1][0]).view(1, 1, 28, 28).to(device)
      im = Variable(im, requires_grad = True)
      digit = Variable(torch.tensor(digit)).to(device).view(1)
      # Put your code here:
      # __begin
      optim = torch.optim.Adam ([im, ], lr = 4e-4, weight_decay=1e-3)

      if lossF == None:
          lossF = nn.CrossEntropyLoss()

      for epoch in range(iters):
          correct = 0

          output = model(im)
          loss = lossF (output, digit)

          loss.backward()

          optim.step()
          optim.zero_grad()


      #__end

      return im
```

```
[40]: ## Let's run our function for 7 and see what image we get.

      im = train_im (net, train_loader, digit=8, iters = 10000)
      im = np.asarray(im.view(28, 28).cpu().detach())
      plt.imshow(im)
```

```
[40]: <matplotlib.image.AxesImage at 0x7f6f012f8e48>
```




In the the following cell we will run the function for all the digits from 0 to 9 and print the outputs!


Enjoy you are done!

```
[36]: for i in range (10):  
      im = train_im (net, train_loader, digit=i, iters=10000)  
      im = np.asarray(im.view(28, 28).cpu().detach())  
      print (i)  
      plt.imshow(im)  
      plt.show()
```


0




1



2




3




4







5




6




7



8




9




[37]: *# Same Procedure as above, but with the better performing model. The features\_*  
*→ learned are not as smooth for the better model imo*

```
for i in range (10):  
    im = train_im (net2, train_loader, digit=i, iters=10000)  
    im = np.asarray(im.view(28, 28).cpu().detach())  
    print (i)  
    plt.imshow(im)  
    plt.show()
```


0




1




2




3




4




5




6



7



8



9



